DTU Compute
Department of Applied Mathematics and Computer Science

Machine Learning for Semi-Automated Annotations of the Linear A Inscriptions

Bachelor Thesis

DTU Compute
Department of Applied Mathematics and Computer Science
Technical University of Denmark

Richard Petersens Plads
Building 324
2800 Kongens Lyngby, Denmark
www.compute.dtu.dk

Preface

This Bachelor's thesis was prepared at the Department of Applied Mathematics and Computer Science at the Technical University of Denmark in fulfillment of the requirements for acquiring a Bachelor of Science in Engineering (BSc Eng) degree in General Engineering (Cyber Systems).

Kongens Lyngby, May 23, 2026

Frederik W. L. Christoffersen

Abstract

This thesis is in the workings... [Write something when finished with all thesis chapters]

Acknowledgements

Frederik W. L. Christoffersen
General Engineering (Cyber Systems), DTU
Author of this thesis

Morten Mørup
Professor, DTU Compute
Main supervisor

Ester Salgarella
Research Fellow, AIAS Aarhus / University of Cambridge
Co-supervisor, creator of SigLA, and winner of the 2026 Michael Ventris Award in Mycenaean Studies

Lukas Esterle
Associate Professor, Aarhus University
Co-supervisor

Generative AI tools were occasionally used as support for brainstorming, linguistic editing, code development, and research during the development of this project.

Thanks to Morten for making this project possible. Thanks to Lukas for his tips and support in developing the fundamental idea for it. And many thanks to Ester for her positive attitude, the resources she provided, and the time she put into this project. I am grateful for the opportunity to have worked with you.

Contents

Preface i

Abstract iii

Acknowledgements v

Contents vii

Acronyms ix

List of Figures xi

List of Tables xiii

1 Introduction 1
- 1.1 Background . . . 1
- 1.2 Related work . . . 3
- 1.3 Problem statement . . . 4
- 1.4 Impact – innovation and application . . . 5
- 1.5 Overview of the thesis . . . 6

2 Data 7
- 2.1 Sources . . . 7
- 2.2 Characteristics and construction . . . 7
- 2.3 Limitations and biases . . . 9

3 Methodology 13
- 3.1 Non-interactive segmentation ability . . . 14
- 3.2 Interaction efficiency . . . 30

3.3 Workflow efficiency 32

4 Results and Discussion **37**
4.1 Non-interactive segmentation performance 37
4.2 Tool-side segmentation performance 46
4.3 Workflow impact 46

5 Conclusion **51**
5.1 Summary of findings per research question 51
5.2 Overall contribution 51
5.3 Future Work 51

Bibliography **53**

A An Appendix **59**
A.1 Non-interactive segmentation results 59

Acronyms

SOTA State of the Art

RQ Research Question

ML Machine Learning

CV Computer Vision

GAN Generative Adversarial Network

HT Haghia Triada Tablet

SAM Segment Anything Model

FATESAM Few-shot Adaptation of Training frEe SAM

GFSAM Graph-based Few-shot Segment Anything Semantically

SA-1B Segment Anything 1 Billion

ViT Vision Transformer

CNN Convolutional Neural Network

NLP Natural Language Processing

LLM Large Language Model

MLP Multilayer Perceptron

Otsu Otsu's method

Gaussian Gaussian adaptive thresholding

CannyFill Canny edge detection with morphological closing

SAM+pts SAM with automatic point prompts

YOLO You Only Look Once

Cefael Collections de l'École française d'Athènes en ligne

GORILA Godart and Olivier, Recueil des inscriptions en linéaire A

PDR Project Definition Report

IPR Intellectual Property Rights

IoU Intersection over Union

Dice Dice-Sørensen coefficient

TPE Tree-structured Parzen Estimator

List of Figures

2.1 Example of a raw inscription image and its corresponding ground truth annotation. 8
2.2 Examples of images with different levels of segmentation difficulty: easy, medium, and hard. For each case, the raw image (left) and the corresponding ground truth annotation (right) are shown. 10
2.3 Examples of images processed with the two different binarization thresholds; 0 (old) and 128 (new). 12

3.1 Hyperparameters for Gaussian adaptive thresholding (Gaussian). 18
3.2 Hyperparameters for Canny edge detection with morphological closing (CannyFill). 20
3.3 Architecture of SAM, from C. Zhang et al. (2023). 21
3.4 Architecture of SAM's mask decoder, from Kirillov et al. (2023). 23
3.5 Architecture of Hiera, adapted from Ryali et al. (2023). . . . 23
3.6 Architecture of SAM2, from Ravi et al. (2024). 24
3.7 Architecture of SAM2's mask decoder, from Ravi et al. (2024). 24
3.8 Architecture of FATESAM, from He et al. (2025). 25
3.9 Hyperparameters for SAM with automatic point prompts (SAM+pts) 26
3.10 A visual example of prompt distribution for SAM with automatic point prompts (SAM+pts), overlayed on HT7a. Green points represent foreground point prompts, while red points represent background point prompts. 27
3.11 Yeah. 33
3.12 Yeah. 34
3.13 Yeah. 35

4.1 Dice Score performance per SAM variant. Split by difficulty. With error bars representing the standard error of the mean. Ordered descendingly from the left by mean performance. . 38
4.2 Dice score comparison of variants of FATESAM with the best performing SAM implementation. 39
4.3 Dice score comparison of all models on the evaluation dataset. 40
4.4 Visual comparison of output masks per model from segmentation on HT7a (easy), shown with Dice scores next to the raw image and ground truth. 42
4.5 QQplot of the residuals of the paired t-test comparing Gaussian and SAM+pts. The residuals appear to be approximately normally distributed, which supports the validity of the paired t-test results. Furthermore, the Shapiro-Wilk test for normality of these residuals yielded a p-value of $0.056 > \alpha = 0.05$, indicating that the null hypothesis of normality could not be rejected. 44
4.6 Dice score performance of the interactive version of SAM2, compared to the classical baselines, plotted against the time taken to use each tool. 47
4.7 Amount of prompts placed per image for the interactive version of SAM2. 48
4.8 Average number of point prompts placed per box prompt per image for the interactive version of SAM2. 49
4.9 Workflow time spent per method, split into tool use 49
4.10 This is how to do it 50

List of Tables

A.1 Best TPE hyperparameter optimization results for CannyFill after 100 trials. 59
A.2 Best TPE hyperparameter optimization results for Gaussian after 100 trials. 59
A.3 Best TPE hyperparameter optimization results for SAM+pts after 100 trials. 60
A.4 Raw Dice score data on easy images. Label column represents the name of the document segmented, while all other columns represent Dice scores obtained for those documents for a given model. 61
A.5 Raw Dice score data on medium images. Label column represents the name of the document segmented, while all other columns represent Dice scores obtained for those documents for a given model. 61
A.6 Raw Dice score data on hard images. Label column represents the name of the document segmented, while all other columns represent Dice scores obtained for those documents for a given model. 62
A.7 Summary Statistics per Model of Dice Score on all images. . 63
A.8 Summary Statistics per Model of Dice Score on easy images. 63
A.9 Summary Statistics per Model of Dice Score on medium images. 63
A.10 Summary Statistics per Model of Dice Score on hard images. 64

CHAPTER 1

Introduction

1.1 Background

When the British archaeologist Arthur Evans in the beginning of the 20th century went to excavate Knossos on Crete, he found a huge amount of clay tablets written in a previously unknown linear script. Some documents were clearly of a newer script, which he called *Linear B*, while others were of an older one, which he named *Linear A*. Some fifty years later, the British amateur archaeologist Michael Ventris finally succeeded in deciphering the newer script Linear B, which he proved to be representing Mycenaean Greek, the earliest attested form of Greek. But the same hypothesis did not hold for the older script Linear A, whose encoded language remained unknown, still yet to be deciphered.

A key obstacle to its decipherment has been the small size of its corpus, lately consisting of only around 2,500 inscriptions (about 1,400 if only including those well-preserved enough to be readable) - significantly smaller than that of its younger counterpart that now spans more than 4,500 inscriptions (which are comparatively longer and better preserved).[1,2] But new inscriptions are still being discovered, with the corpus most recently extended in 2025 by over 100 new documents.[3,4]

Another challenge has been the lack of digital representations of these inscriptions, in a form suitable for statistical analysis. While photographs with accompanying drawings of nearly the entire corpus are available online through Collections de l'École française d'Athènes en ligne (Cefael) (see Godart and Olivier (1976-1985)), these are only available in print

[1] Salgarella 2020.
[2] Salgarella 2025, p. 13, 46.
[3] Del Freo and Zurbach 2025.
[4] Consiglio Nazionale delle Ricerche 2025.

form. This motivated the development of the palæographical[5] database *The Signs of Linear A* (SigLA), which seeks to "[develop] a systematic, exhaustive and user-friendly open access database of all Linear A inscriptions [...] in order to carry out statistical and palæographic analyses within the epigraphic corpus" (Salgarella and Castellan 2020). However, the last addition to SigLA was in 2021, and the project has so far managed to document just 772 of all currently known inscriptions.[6,7]

1.1.1 A problem of manual digitization

The drawings currently stored on SigLA were made by Ester Salgarella using the painting tool Krita. For each document, she manually created her own annotated drawings based on the originals stored by Cefael. From these, she proceeded to create layered Krita files with metadata for each sign, which could then at last be imported to SigLA.[8] According to her (personal communication), this has been a highly time-consuming undertaking. And the necessity of this kind of labor-intensive work has arguably been the biggest reason to why SigLA still does not contain the entire available corpus.

The aim of this project was therefore to investigate whether parts of this process can be automated. More specifically, whether some kind of semi-automatic segmentation tool (given just the raw tablet images as input) can accelerate the workflow, while maintaining high-quality outputs that are still compatible with SigLA's database.

1.1.2 A problem of small data

The small size of the Linear A corpus and the irregularities in quality of its documents is a huge problem for the application of deep learning approaches, which typically require large amounts of data to perform well. Consequently, this project instead sought to make use of generalized Machine Learning (ML) and Computer Vision (CV). In partic-

[5]The term "palæographical" refers to the study of ancient writing systems.
[6]SigLA 2021a.
[7]SigLA 2021b.
[8]Salgarella and Castellan 2020.

ular Meta's State of the Art (SOTA) interactive segmentation model, the Segment Anything Model (SAM) (Kirillov et al. 2023), which is pre-trained on a massive and diverse dataset - Segment Anything 1 Billion (SA-1B), of over 1 billion masks. SAM's generalization and interactivity (taking user-placed prompts that guide the segmentation) potentially allows for semi-automatic annotation of Linear A inscriptions without the need for large amounts of training data, which was unfortunately not available in this context. It was therefore of particular interest to map the capabilities and limitations of SAM within this domain.

1.2 Related work

Within the last decade, the field of document annotation has seen significant advancements through the use of machine learning. Take for example the EU-funded generic CNN-based framework for historical document processing *dhSegment* (Oliveira, Seguin, and Kaplan 2018). Or the more recent work on segmentation of stone inscriptions (P. Zhang, C. Li, and Sun 2024), where SAM and other state-of-the-art methods are beaten as benchmarks by their stacked U-Net and GAN architecture. Or even the YOLO and Roboflow based approach to segmentation of Palmyrene Aramaic inscriptions (Hamplová et al. 2024) that augment[9] their dataset with a Generative Adversarial Network (GAN) to increase performance.

Nevertheless, these methods all depend on deep learning, making them inconvenient for domains of small data, which instead seem to require a generalized and interactive approach such as SAM.

1.2.1 SAM and its applications

Ever since its release in 2023, SAM has been extensively evaluated across a wide range of CV tasks to investigate whether it can really "segment anything". Papers have evaluated its performance within niche domains, such as medical images (Ma et al. January 2024), camouflaged objects (Tang, Xiao, and B. Li 2023), transparent objects (Han et al. 2023), hierarchical texts (Ye et al. 2024), and even oracle bone inscriptions (Wang

[9]The term *augmenting* refers to the process of increasing the size of a dataset.

et al. 2025). Most of these papers found that SAM still tends to struggle with such domains, despite its impressive zero-shot[10] generalization. Meanwhile, some of them also found that this performance can be significantly improved through architectural extensions and domain-specific fine-tuning, which unfortunately requires deep learning and a lot of data. However, equally, it has been shown that SAM's performance can significantly improve through effective prompt engineering alone (Price, Rundensteiner, and Cote 2025), requiring no deep learning at all. And within the domain of 3D medical images, a training-free few-shot adaptation of SAM2 (see Ravi et al. 2024) has been recently demonstrated to surpass supervised methods like U-Net, and even the best fined-tuned versions of SAM (He et al. 2025). Thus, it seemed reasonable to test out the limits of SAM's generalization and interactivity within the domain of Linear A document images, to investigate whether it could be used to accelerate the annotation workflow of Ester Salgarella.

1.3 Problem statement

As previously stated, the aim of this project was to investigate whether parts of the process of creating annotated drawings like the ones in SigLA can be automated. More specifically, whether some kind of semi-automatic segmentation tool using interactive and generalized CV - specifically the Segment Anything Model (SAM) - can accelerate the workflow, while maintaining high-quality outputs that are still compatible with a database like SigLA. But what would be the relative segmentation quality of such a tool? How automatic would it be exactly? And by how much could it improve the efficiency of the workflow? To concretize the research, the problem statement was decomposed into three Research Questions (RQs) - each accompanied by their respective operationalization (see the following subsection for details).

[10]The term *zeroshot* refers to model performance on hitherto unseen data.

1.3.1 Research questions

1. **RQ1 (Segmentation quality)**: *How well does a SAM-based segmentation approach perform quantitatively in terms of segmentation accuracy, compared to simple classical image processing baselines when applied to Linear A document images?*

   Operationalization: SAM-based masks were compared to classical image-processing baselines using overlap metrics (i.e. Dice) against the ground truth segmentations derived from SigLA's Krita files using a self-made dataset of 60 documents.

2. **RQ2 (Interaction efficiency)**: *To which extent can user interaction be reduced by a SAM-based semi-automatic segmentation workflow, while still maintaining high-quality outputs?*

   Operationalization: Two self-developed segmentation tools - one SAM-based, one classical - were evaluated in a case study of 3 documents. The number of Ester Salgarella's tool-side user interactions (i.e. prompts placed for SAM, and parameter adjustments for the classical method) was compared against segmentation quality (i.e. Dice) relative to ground truth (SigLA's Krita files).

3. **RQ3 (Workflow efficiency)**: *How much reduction in annotation time can be achieved using the proposed tool compared to fully manual digitization, under controlled experimental conditions?*

   Operationalization: The time it took Ester Salgarella to annotate manually compared to being assisted by two self-developed segmentation tools - one based on SAM, and the other based on a classical baseline - was measured in a case study of 3 documents.

1.4 Impact – innovation and application

The work of this thesis presents itself as a first proof-of-concept (and a mapping of the difficulties involved) in using SAM and modern CV to accelerate the annotation of Linear A documents. Its main contribution

is a prototype of a semi-automatic annotation tool that extracts engravings from raw tablet images, along with an evaluation framework that can be used to assess the performance of that tool and any potential future derivatives. Such a tool could potentially also be applied to Linear B, as it is the whole field of Aegean studies that seems to suffer from severe under-digitization of the primary evidence (which is a huge problem for the scientific investigations of scholars). Hopefully, the work will inspire further research and development within the field of digital epigraphy. Particularly in the application of generalized and interactive CV approaches such as SAM on Linear A, and perhaps even in the broader context of use on small corpus ancient writing systems, where the lack of big data is ill fit for approaches using deep learning.

1.5 Overview of the thesis

The rest of this thesis is structured into four main chapters (Data, Methodology, Results and Discussion, Conclusion), all split up into sections concerning each of the three research questions.

CHAPTER 2

Data

2.1 Sources

The dataset used in this thesis was constructed from 60 images of raw Linear A inscriptions, and their corresponding ground truth annotations. Cefael (see Godart and Olivier 1976-1985) was chosen as the source for raw document images, since it is the only online and widely publicly available source of the printed five volume corpus - Godart and Olivier, Recueil des inscriptions en linéaire A (GORILA). As ground truth, annotations provided by Ester Salgarella from SigLA (see Salgarella and Castellan 2020) were chosen instead of the annotated drawings from the corpus, as the former are more up to date, and of better quality.

What follows in this chapter is a detailed description of this dataset, documenting its characteristics, its construction process, and its limitations and biases.

2.2 Characteristics and construction

The dataset was made to consists of 60 chosen image pairs of raw documents and their corresponding ground truth annotations. Each of these image pairs was categorized into one of three difficulty levels (easy, medium, hard) based on the visual quality of the inscriptions (see Section 2.2.1). They were then split into a test set of 27 images, a validation set of 27 images, and a support set of 6 images, with each set containing an equal proportion of images from each of the three difficulty levels.

Only inscriptions with clay as writing support were included, as they are the most common and representative of the Linear A corpus, and were the only ones for which SigLA ground truth annotations were al-

ready available. The raw document images from Cefael are screenshots of grayscale images placed within the scanned printed edition of the corpus, and are thus not of the highest quality, being the only widely publicly available images of the Linear A inscriptions. For an example of a raw document image and its corresponding ground truth annotation, see Figure 2.1.

((a)) Raw image (HT7b) **((b))** Ground truth (HT7b)

Figure 2.1: Example of a raw inscription image and its corresponding ground truth annotation.

2.2.1 Difficulty classification

The dataset was categorized into three difficulty levels (easy, medium, hard) based on the visual quality of the raw images and the clarity of the inscriptions. This classification was performed manually by inspecting each image and considering factors such as contrast, noise, and the presence of tablet breakages or general damage to the inscriptions. The concrete criteria used to classify the difficulty sets were as follows:

- **Easy:** labeled as such for documents with clear ingravings, good contrast, minimal noise, and constant lighting.
- **Hard:** labeled to those with poor contrast, significant noise, and substantial damage making it hard even for the naked eye to read the inscriptions.
- **Medium:** labeled to all those inbetween the qualities of easy and hard.

For visual examples of images within each of these difficulties, see Figure 2.2.

2.2.2 Image registration and preprocessing

To obtain any kind of meaningful comparison between model outputs and ground truth, the raw images and their ground truth annotations had to be aligned and registered to each other. This was done manually in Krita for all 60 image pairs by downscaling the ground truth images to the individual scales of their raw counterparts, and aligning by translation and rotation, using the raw images as background reference.

2.3 Limitations and biases

2.3.1 Quality of the raw tablet images

Due to the weak quality of the raw tablet images, they carry only parts of the visual information of the tablets they represent. When the drawings of GORILA were made, the researchers had access to the tablets themselves, meaning they could see what the images did not pick up, and better discern actual engravings from mere damages to the tablets. Therefore, these tablet images are far from the best source of data there is, being however, the only source publicly available. That lack of good quality data can unfortunately limits the output quality of segmentations constructed therefrom.

((a)) Easy KHWc2104 (raw)

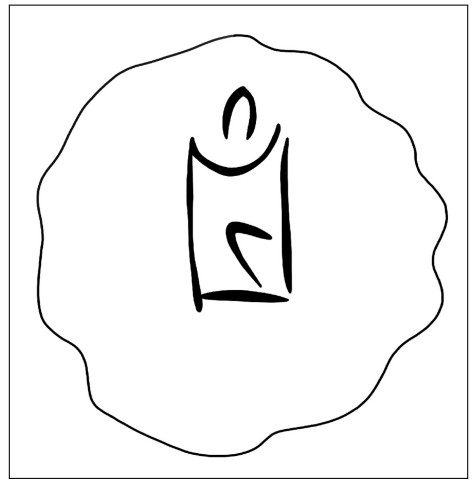

((b)) Easy KHWc2104 (gt)

((c)) Medium MA1a (raw)

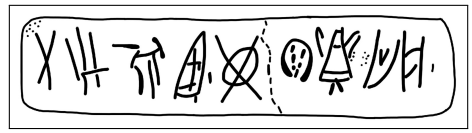

((d)) Medium MA1a (gt)

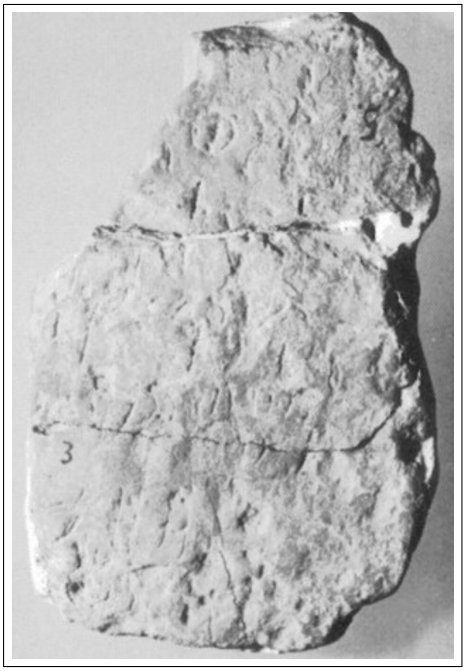

((e)) Hard HT3 (raw)

((f)) Hard HT3 (ground truth)

Figure 2.2: Examples of images with different levels of segmentation difficulty: easy, medium, and hard. For each case, the raw image (left) and the corresponding ground truth annotation (right) are shown.

2.3.2 Alignment to the ground truth

Even after image registration, ground truth images still did not perfectly align with the raw images, affecting the obtained segmentation quality. This misalignment was a source of noise in the evaluation. Ester's annotations were never constructed by overlaying them on the raw images, but rather by doing so on the accompanied corpus drawings, only using the raw images as reference. Thus, the raw images and the ground truth annotations were never perfectly aligned to begin with, and the manual registration process could only do so much to mitigate that. This is seen by the low Dice values obtained for all models in the results (see Section 4.1).

2.3.3 Binarization of the drawings

Due to a missing threshold step of the registered JPEGs in the initial registration of the dataset, the ground truth annotations were to begin with binarized with a threshold of 0 instead of 128, meaning that only the pure black pixels of the drawings were considered as part of the ground truth foreground, while the lighter gray pixels were not accounted for. This led to lower Dice scores in general, as the ground truth in practice was thinner and more sparse compared to the actual drawings.

This issue was for the evaluation later corrected to a threshold of 128, meaning that all pixels darker than 128 now were considered foreground, and all pixels lighter than that were considered background (see Figure 2.3 for visual examples of the two different ways of binarization). However, only the hyperparameter optimization of the models was not redone for the new binarization format, and was thus performed with ground truth in the old binarization format.

Arguably, this only affected the effectiveness of the optimization and what the models were optimized for (penalizing non-conservative segmentation styles more than otherwise), while still preserving fairness and generalizability, as all model hyperparameters were tuned under the same conditions.

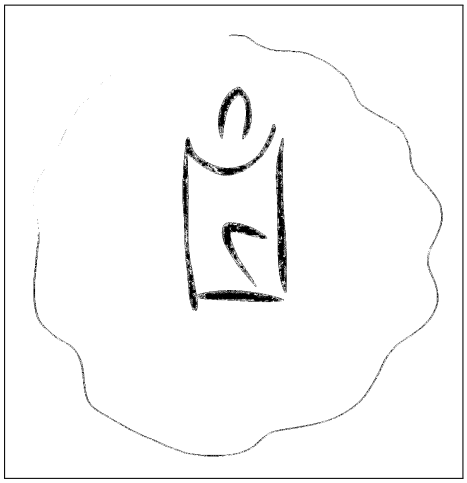

((a)) KHWc2104 (old)

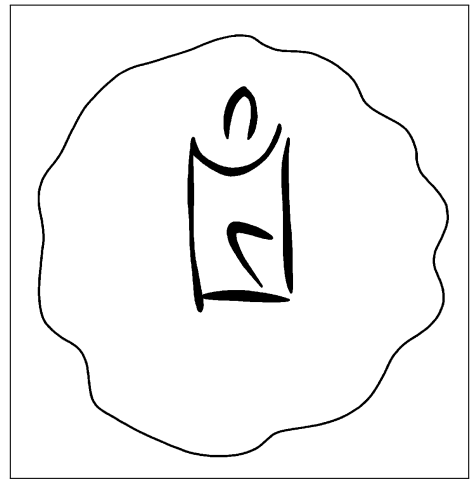

((b)) KHWc2104 (new)

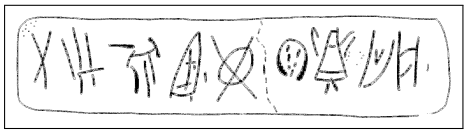

((c)) MA1a (old)

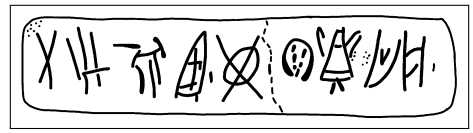

((d)) MA1a (new)

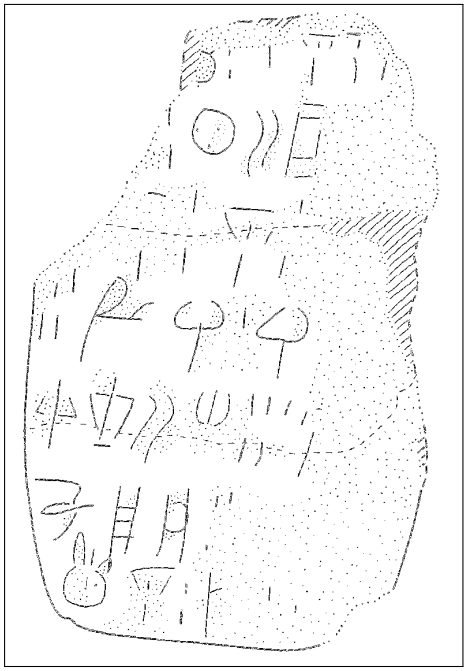

((e)) HT3 (old)

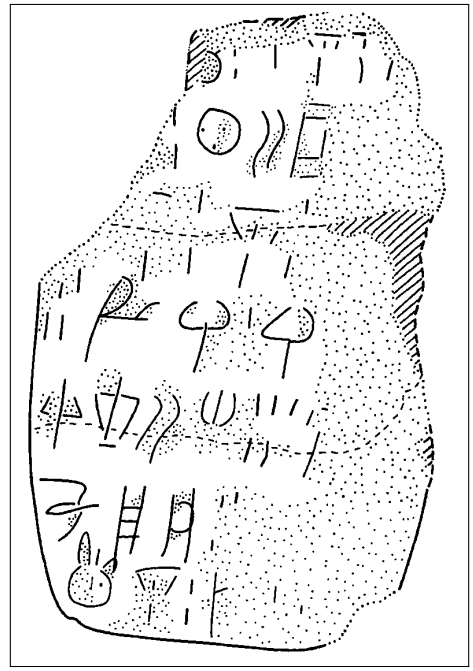

((f)) HT3 (new)

Figure 2.3: Examples of images processed with the two different binarization thresholds; 0 (old) and 128 (new).

CHAPTER 3
Methodology

This chapter describes in detail the operationalization of answering the question of whether a SAM-based segmentation tool can accelerate the annotation workflow of Linear A documents (see Section 1.3.1 for an overview of the individual RQs). Overall, this operationalization was achieved through a comprehensive evaluation of segmentation quality (see Section 3.1), and a case study involving Dr. Ester Salgarella measuring the interaction efficiency (see Section 3.2) and workflow efficiency (see Section 3.3) of a self-developed SAM-based segmentation tool. In that way, these three abstract virtues were decomposed into independent measurable metrics, making it possible to concretely answer the research questions, and where the potential benefits and drawbacks of SAM lie.

3.0.1 Development environment and tools

All software written for this thesis was made publicly available on GitHub (Christoffersen 2026). Python 3.13 was used for all code development except frontend. The choice of this programming language was made due to its big community and library support for nearly everything; be it plotting (Matplotlib), matrix math (NumPy), image processing (OpenCV), etc. Python's package manager (pip) also made it exceptionally easy to install and manage dependencies, and set up the development environment. For the ML and SAM implementation libraries needed, Python has also been the language of choice, and thus it was the obvious choice for development. The particular version was chosen, as it was the second latest stable release at the time of development, with the risk of Python 3.14 still not having full support by all libraries.

3.1 Non-interactive segmentation ability

To answer the question of RQ1: *How does the initial segmentation quality of SAM compare to that of classic segmentation methods?*, the initial fully-automatic segmentation quality of SAM was evaluated against a set of classical baseline models, using a suitable segmentation metric (i.e. Dice score) computed from the pairs of predicted and ground truth masks. All models were tuned by hyperparameter optimization to their best performance on the dataset, and the best performing version of SAM was compared against the best performing baseline model.

The following subsections therefore document the configuration of the baseline and state-of-the-art models, as well as the evaluation protocol used to answer RQ1.

3.1.1 The problem in mathematical terms

The segmentation process can be mathematically described as a function that - given some predicate function f, and an input image I - produces a binary mask B of the same dimensions, where each pixel value indicates whether it belongs to the foreground (here denoted as 1) or the background (denoted as 0):

$$B(x,y) = \begin{cases} 1 & \text{if } f(I,x,y) \\ 0 & \text{otherwise} \end{cases}$$

In a threshold-based approach for example, the predicate function can be expressed as $f(I,x,y) = I(x,y) > T$, where T is the threshold value that determines the cutoff between foreground and background pixels.[1]

[1] Gonzalez and Woods 2018, p. 743.

3.1.1.1 Convolutional kernels

[2]

$$(w \star I)(x, y) = \sum_{i=-a}^{a} \sum_{j=-b}^{b} w(i, j) I(x - i, y - j)$$

where w is the convolution kernel (window) of size $m \times n$, and $a = (m-1)/2$ and $b = (n-1)/2$ are the half-width and half-height (assuming odd dimensions), respectively.[3]

3.1.2 Baseline model configuration

The baseline models against which the best SAM implementation was evaluated (from simplest to most advanced) were:

- Global thresholding (Otsu's method)
- Local adaptive thresholding (Gaussian)
- Edge-based segmentation (Canny edge detection with region filling)

Together, they cover a range of classic segmentation techniques that might be used for binary foreground extraction. What follows is a brief description of each of these methods, including their characteristics, how they work, and their implementation details.

3.1.2.1 Bilateral pre-filtering

As a pre-filtering step for all the classical models, a bilateral filter (see Tomasi and Manduchi 1998) was applied to the input image before feeding it to the classical methods. In this way, more noise could be reduced while preserving engravings, as this filter is known to be effective at reducing noise while preserving edges. It takes three hyperparameters: the kernel window diameter d, the spatial standard deviation σ_{space}, and the intensity standard deviation σ_{color}. Discretely, its output intensity $I_{\text{bf}}(p)$

[2] Gonzalez and Woods 2018, p. 159.
[3] Gonzalez and Woods 2018, p. 157.

at a given pixel position p (where q is the position of one of its neighbors within the kernel window) can be defined as follows:[4]

$$I_{\text{bf}}(p) = \frac{\sum_q I(q)\, W_{\text{space}}(p,q)\, W_{\text{color}}(p,q)}{\sum_q W_{\text{space}}(p,q)\, W_{\text{color}}(p,q)}$$

Here, the spatial weight based on the distance between pixels p and q, and the color range weight based on the intensity difference between p and q are respectively given by:

$$W_{\text{space}}(p,q) = \exp(-\frac{\|p-q\|^2}{2\sigma^2_{\text{space}}})$$

$$W_{\text{color}}(p,q) = \exp(-\frac{\|I(p)-I(q)\|^2}{2\sigma^2_{\text{color}}})$$

It was decided to use the same σ value for both spatial and color weights, as this simplified the hyperparameter tuning (see Section 3.1.4.2) without significantly compromising the filter's performance.

3.1.2.2 Global thresholding (Otsu's method)

Otsu's method is a global thresholding method that works on gray-level histograms, which was originally proposed by Otsu (1979). As the simplest of the baselines, it represents the minimum acceptable segmentation quality, since it only takes the global gray level distribution into account, with no assumptions about any spatial relationships between pixels. It is therefore expected to perform poorly on images with erosion, uneven lighting, and differences in surface texture[5], which are the exact characteristics of the images to be segmented.

Here is how it works: First, let pixels of the raw image be divided into L shades of gray $[1, 2, ..., L]$. Then, let each pixel be assigned one of two classes C_0 or C_1, guided by a threshold T, where C_0 contains the shades $[1, 2, ..., T]$, and C_1 contains the rest $[T+1, T+2, ..., L]$. Otsu's method then works by performing an exhaustive search over all L gray levels to

[4] Pan 2025.

[5] Gonzalez and Woods 2018, p. 751.

find an optimal threshold T^* that results in a maximum between-class variance σ_B^2:

$$T^* = \underset{1 \leq T < L}{\arg\max}\ \sigma_B^2(T)$$

The between-class variance σ_B^2 is then defined by a function of the respective total probabilities of the two classes (ω_0 and ω_1) and the mean gray level value per class (μ_0 and μ_1):

$$\sigma_B^2(T) = \omega_0(T)\omega_1(T)[\mu_1(T) - \mu_0(T)]^2$$

3.1.2.3 Local adaptive thresholding (Gaussian)

Local adaptive thresholding works by computing a threshold for each pixel based on its local neighborhood. This is an advantage in cases where the image has varying light conditions (which is the case for the images to be segmented), as the threshold adapts to local lighting.

Mathematically, the local threshold T is a function of a given pixel position (x, y) computed by a weighted sum of the neighborhood. The generated binary mask B is then defined as:[6]

$$B(x, y) = \begin{cases} 1 & \text{if } I(x, y) > T(x, y) \\ 0 & \text{otherwise} \end{cases}$$

To compute $T(x, y)$, there are a few weighting functions to choose from. For this baseline, the Gaussian kernel was used, as it can limit the influence of distant noisy pixels by assigning weights based on distance from the kernel center. With a Gaussian kernel G_σ, the threshold $T(x, y)$ is computed as a Gaussian-weighted sum of local intensities subtracted by a constant C. Using convolution ($\star$), this can be expressed as:[7]

$$T(x, y) = (G_\sigma \star I)(x, y) - C$$

In general, the weight at an offset (i, j) from the kernel center can be written as follows (where α is a normalization constant ensuring that

[6] Gonzalez and Woods 2018, p. 761.
[7] OpenCV 2024.

the kernel sums to 1, and σ controls the spatial spread of the kernel):[8]

$$G_\sigma(i,j) = \alpha \cdot \exp\left(-\frac{i^2+j^2}{2\sigma^2}\right)$$

This gave in total 4 hyperparameters for this baseline (see Figure 3.1 for a list of these):

X

1. C: the constant subtracted from the Gaussian-weighted sum to compute the local threshold (see Section 3.1.2.3),
2. $d_{gaussian}$: the diameter of the kernel window for Gaussian adaptive thresholding (see Section 3.1.2.3),
3. $d_{bilateral}$: the diameter of the kernel window for bilateral prefiltering (see Section 3.1.2.3),
4. σ: the spatial and intensity standard deviation for bilateral prefiltering (see Section 3.1.2.3),

Figure 3.1: Hyperparameters for Gaussian adaptive thresholding (Gaussian).

[8] Gonzalez and Woods 2018, p. 167.

3.1.2.4 Edge-based segmentation (Canny edge detection with region filling)

Canny edge detection is a method that by itself only segments the edges of an image, excluding the enclosed regions within those edges. Here is how it works: First, the input image is smoothed using a Gaussian filter to reduce noise (using the Gaussian kernel as defined in the previous section):

$$I_s(x, y) = (G_\sigma \star I)(x, y)$$

Then, for the image gradient of each pixel, the magnitude $M_s(x, y)$ and the orientation $\theta_s(x, y)$ are computed:[9]

$$M_s(x, y) = \| \nabla I_s(x, y) \| = \sqrt{\left(\frac{\partial I_s}{\partial x}\right)^2 + \left(\frac{\partial I_s}{\partial y}\right)^2}$$

$$\theta_s(x, y) = \arctan\left(\frac{\partial I_s / \partial y}{\partial I_s / \partial x}\right)$$

Next, *non-maxima suppression*[10] is applied to retain only local maxima of $M_s(x, y)$ along the gradient direction $\theta_s(x, y)$, resulting in thinner edges. At last, this is followed by *hysteresis thresholding*[11], classifying pixels as strong, weak, or non-edges with a low threshold T_l and a high threshold T_h, only preserving weak edges if connected to strong edges.

Inspired by (Rosebrock 2015), these thresholds were automatically determined by setting

$$T_l = \max\Big(0, (1 - \sigma)\, v\Big), \quad T_h = \min\Big(255, (1 + \sigma)\, v\Big)$$

turning σ into a hyperparameter that together with the grayscale median of the given image v controlled the sensitivity of the edge detection.

In an attempt to fill the regions enclosed by the detected edges, morphological closing (dilation followed by erosion) was applied to the edge mask. This gave in total 2 hyperparameters (see Figure 3.2 for a list of these).

[9] Gonzalez and Woods 2018, p. 730.
[10] Gonzalez and Woods 2018, p. 730-731.
[11] Gonzalez and Woods 2018, p. 732.

X

1. σ: the constant used to determine the thresholds T_l and T_h for the hysteresis thresholding,

2. $d_{closing}$: the diameter of the kernel window for the morphological closing of the edges

Figure 3.2: Hyperparameters for Canny edge detection with morphological closing (CannyFill).

3.1.3 State-of-the-art

To be compared against the classical baselines was the generalized deep learning Segment Anything Model (SAM) (see Kirillov et al. 2023). Due to its popularity as a SOTA segmentation model, and the fact that its different use cases require different trade-offs in terms of accuracy, speed, and domain adaptation, there are many different variants of SAM.

For the sake of fairness, it seemed only reasonable to test out a range of them, including of course the original, the computationally efficient MobileSAM (see Zhao et al. 2023), and the newer SAM2 (see Ravi et al. 2024). As the checkpoints (weight files) needed to run the newly released SAM3 (Carion et al. 2025) were still access-restricted at the time of development, it was not included in the evaluation. However, few-shot implementations of SAM - like Few-shot Adaptation of Training frEe SAM (FATESAM), and Graph-based Few-shot Segment Anything Semantically (GFSAM) - were also included in the evaluation, because of their potential to better accommodate to niche domains. To be able to run these computationally demanding variants of SAM, model execution was outsourced to Modal (see Modal Labs 2026) - a serverless cloud computing platform usually used for ML inference, tuning and training.

The following subsections describe the architecture of the chosen variants, and how they were implemented for the evaluation.

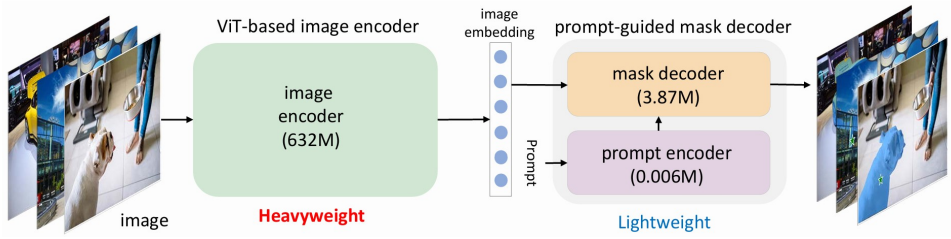

Figure 3.3: Architecture of SAM, from C. Zhang et al. (2023).

3.1.3.1 SAM architecture

For the original SAM and MobileSAM, the model architecture consists of an image encoder, and a prompt-guided mask decoder (see Figure 3.3). The image encoder used is a Vision Transformer (ViT)[12] whose purpose is similar to that of a Convolutional Neural Network (CNN), processing an input image to extract its features. But instead of convolutional layers, a ViT makes use of a transformer-based encoder traditionally used in Natural Language Processing (NLP) and Large Language Models (LLMs). As an alternative to words, it treats an image as a sequence of fixed size patches (regions of the image) embedded into vector representations with positional encodings, similar to how tokens work in NLP. These embedded patches are then processed through a repeated series of normalization, feed-forward Multilayer Perceptrons (MLPs), and in-between those, a mechanism known as self-attention that weights the patches according to their relevance to eachother.[13]

More specifically, given the input embedding X, self-attention computes the projections; quieries: Q, keys: K, and values: V, as products of X with respective learned matrices, and passes the information on as a new matrix (where d_k is the dimension of Q and K):

$$Attention(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

[12]The reason that MobileSAM is computationally faster than the original is its lighter ViT implementation TinyViT (Wu et al. 2022) (with only 5M parameters) compared to those of SAM (ViT-H, ViT-L, ViT-B with 632M, 307M, and 86M, respectively).

[13]see Vaswani et al. 2023, for the math behind *attention*.

Here, the softmax function normalizes the output to a probability distribution (putting each element on the interval $]0, 1[$, with their sum being 1), which is then multiplied by V (again turned into logits / non-normalized probabilities).

Thanks to this attention mechanism, after a number of these series-layers, the ViT finally outputs an image embedding containing the spatial context of the image.[14] It is this image embedding, combined with embedded prompts from the prompt encoder (referred to as prompt tokens) that are fed to the mask decoder to generate the final segmentation mask. However, before this happens, the prompt encoder takes any given sparse prompts (i.e. box or point prompts) as embeddings by summing their positional encodings with learned type-specifc embeddings (one foreground/background embedding per point prompt, and two opposite corner embeddings per box prompt).

These are then fed to the lightweight mask decoder (which is also transformer-based) along with the image-embedding from the ViT, and 3 learnable output token embeddings (see Figure 3.4). Self-attention is applied to the prompt and output tokens, and two-way cross-attention (same as self-attention but with one embedding as queries, and the other as keys and values, and vice-versa) is applied between the tokens and image embedding for them to influence and attend one another. To keep track of spatial coordinates and prompts, positional encodings and the original prompt embeddings are added as input at every attention layer.

At the end of this process, the image embedding is upsampled by convolution to the original image resolution. The three output tokens are then processed in parallel by two separate MLPs: one that generates their final segmentation masks by a dot product with the output and the upscaled image embedding, and another one that outputs a predicted IoU score for each mask. For single output masks, the mask is generated from a fourth output token.[15]

[14] see Dosovitskiy et al. 2021, for how ViTs work in general.
[15] Kirillov et al. 2023, p. 17.

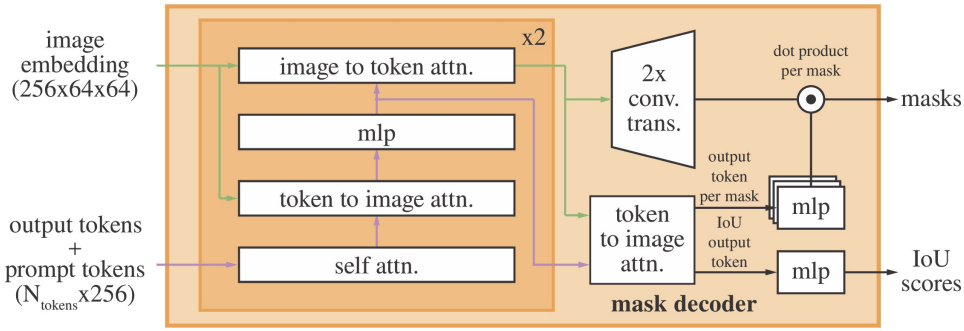

Figure 3.4: Architecture of SAM's mask decoder, from Kirillov et al. (2023).

3.1.3.2 SAM2

With the release of SAM2 (Ravi et al. 2024), the image encoder was changed to Hiera (Ryali et al. 2023) - a hierarchical ViT, which is more parameter efficient than vanilla ViTs, as tokens are pooled together at different stages of the encoding process, lowering attention complexity (see Figure 3.5). Additionally, a memory attention mechanism was added for consistency, allowing the model to attend to segmentation outputs of previous frames, making it easier to segment videos.

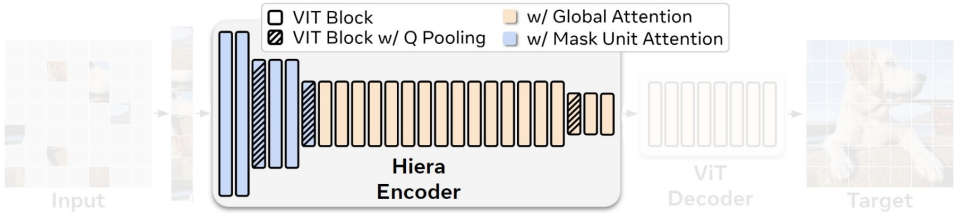

Figure 3.5: Architecture of Hiera, adapted from Ryali et al. (2023).

3.1.3.3 FATESAM

Within the domain of 3D medical images, FATESAM - a training-free few-shot adaptation of SAM2 (see He et al. 2025) has been recently

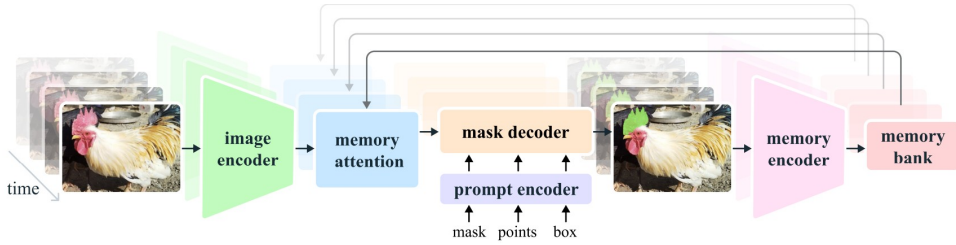

Figure 3.6: Architecture of SAM2, from Ravi et al. (2024).

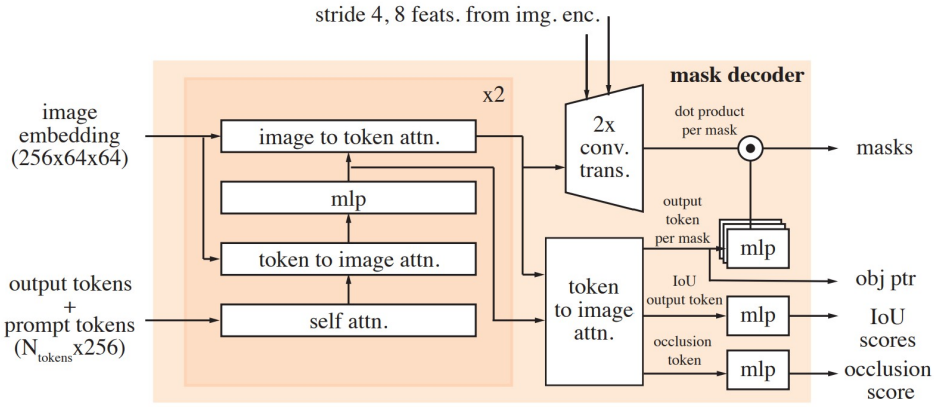

Figure 3.7: Architecture of SAM2's mask decoder, from Ravi et al. (2024).

demonstrated to surpass even the best fined-tuned versions of SAM. It works in principle by making use of SAM2's memory attention, by passing the few-shot support images to its memory encoder (see Figure 3.8), giving the mask decoder context to generate similar segmentation masks for the input query image. The support images are chosen by a feature similarity ranking to the input to make sure they are as representative as possible.

To adapt FATESAM to single frame images, the object pointers were turned off through Hydra/YAML overrides to avoid shape mismatches.

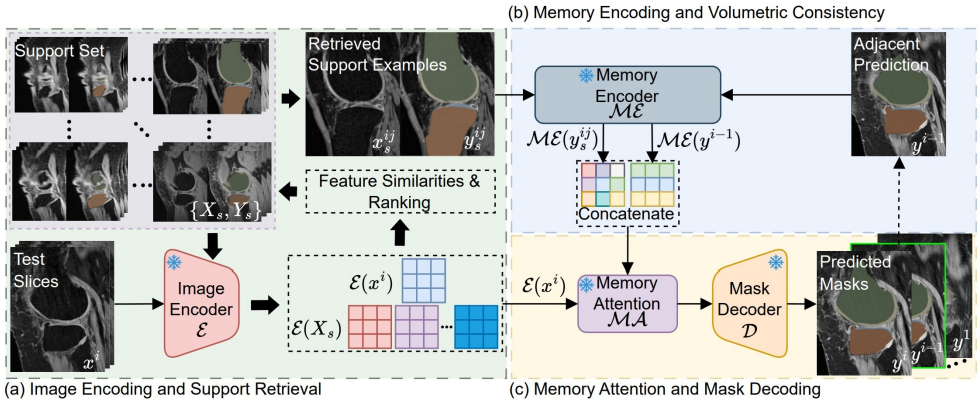

Figure 3.8: Architecture of FATESAM, from He et al. (2025).

3.1.3.4 GFSAM

3.1.3.5 Threshold-inferred automatic point prompts

Another way to adapt SAM to the specific characteristics of the domain is through the use of automatic prompts. But as it is particularly difficult to automatically create meaningful box prompts without a trained object detection model, only point prompts were used for this adaptation.
Inspired by point of interest placement in (Price, Rundensteiner, and Cote 2025), automatic point prompts for SAM (see Figure 3.10 for a visual example) were derived from the output mask of the classical Gaussian adaptive thresholding baseline (see Section 3.1.2.3), This was done by randomly sampling a number of foreground and background pixels from the mask as prompts for SAM+pts. Concretely, the foreground cloud of prompts were sampled from the foreground of the mask by n_{fgd} points, while the background cloud of prompts were sampled from the erosion of the inverse of the mask by n_{bgd} points. This gave in total 7 hyperparameters (see **??**) for SAM+pts.

X

1. C: the constant subtracted from the Gaussian-weighted sum to compute the local threshold (see Section 3.1.2.3),
2. $d_{gaussian}$: the diameter of the kernel window for Gaussian adaptive thresholding (see Section 3.1.2.3),
3. $d_{bilateral}$: the diameter of the kernel window for bilateral prefiltering (see Section 3.1.2.3),
4. σ: the spatial and intensity standard deviation for bilateral prefiltering (see Section 3.1.2.3),
5. n_{fgd_points}: the number of foreground point prompts to be sampled from the foreground of the first Gaussian thresholded mask,
6. n_{bgd_points}: the number of background point prompts to be sampled from the sure background region,
7. $d_{gap\_erosion}$: the kernel size for erosion used to create a gap between the foreground and background regions.

Figure 3.9: Hyperparameters for SAM with automatic point prompts (SAM+pts)

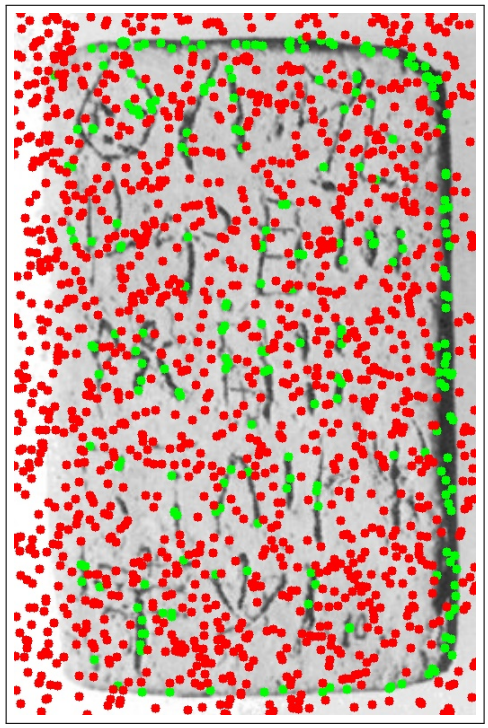

Figure 3.10: A visual example of prompt distribution for SAM with automatic point prompts (SAM+pts), overlayed on HT7a. Green points represent foreground point prompts, while red points represent background point prompts.

3.1.4 Evaluation protocol and performance metrics

For an evaluation of initial segmentation quality, relevant metrics and statistical tests needed to be defined to determine whether one model performed significantly better than another. Furthermore, all models with hyperparameters needed to be optimized to their best performance on the dataset, to ensure a fair comparison. Naturally, due to the nature of RQ1, it was of particular interest to compare the best performing baseline model against the best performing version of Segment Anything Model.

What follows is therefore a description of the choice of metrics, the optimization of model hyperparameters, and the statistical tests for the evaluation, used to answer RQ1: *How does the initial segmentation quality of SAM compare to that of classic segmentation methods?*.

3.1.4.1 Choice of metrics

When it comes to evaluation of segmentation models, the two most commonly used metrics are the Dice-Sørensen coefficient (Dice) and the Intersection over Union (IoU). The Dice score is most commonly used in cases where there are heavy class imbalances (e.g. lots of background compared to regions segmented), and when foreground objects are small and thin. It is naturally more foregiving than the related segmentation metric IoU, which instead tends to be rather strict, and is specifically used to evaluate segmentation results where small displacements are critical for results. Therefore, segmentation performance was reported using Dice. Mathematically, it is defined as twice the size of the intersection of the predicted segmentation mask $\hat{Y}$ and the ground truth mask Y, divided by their summed size:[16]

$$\text{Dice}(\hat{Y}, Y) = \frac{2 \cdot |\hat{Y} \cap Y|}{|\hat{Y}| + |Y|}$$

[16] Mohammadi, Mollazade, and Behroozi-Khazaei 2025.

Thus, the raw metrics ended up consisting of the Dice score for each image, and each model.

3.1.4.2 Hyperparameter optimization

To ensure a fair comparison between the models, all models with hyperparameters were tuned to their best performance on the dataset using Optuna, a hyperparameter optimization framework. The optimization itself was performed using the Tree-structured Parzen Estimator (TPE) algorithm (see Bergstra, Yamins, and Cox 2013)[17], which is a Bayesian optimization method that models the relationship between hyperparameters and performance to efficiently explore the hyperparameter space and find the optimal ones. Hence, it is more sophisticated than a random search, which just samples different hyperparameter combinations without considering the results of past evaluations.

It works by building two probabilistic models: one for the hyperparameter configurations that have resulted in good performance, and another for those that have resulted in poor performance. Then, it uses these models to select the next hyperparameter configuration to evaluate, by maximizing the expected improvement over the current best performance. This way, it can focus the search on promising regions of the space of possible hyperparameters without having to rely on chance or exhaustive search.

3.1.4.3 Statistical tests

To determine whether the observed differences in scores between models were statistically significant, paired t-tests were conducted. The paired t-test has two main assumptions: 1. independence of observations, and 2. normality of residuals. The first assumption is already satisfied a priori, since the segmentation performance of one image does not directly influence the performance on another. However, the second assumption depends on results. Thus, the residuals of the paired differences in Dice scores between the two models in question needed to be checked for normality using a QQ-plot and the Shapiro-Wilk test. And in case the

[17] also see Bergstra, Bardenet, and Kégl December 2011.

normality assumption was violated, the non-parametric domain-specific Wilcoxon signed-rank test would be used instead.[18]

3.2 Interaction efficiency

To answer the question of RQ2: *To which extent can user interaction be reduced by a SAM-based semi-automatic segmentation workflow, while still maintaining high-quality outputs?*, a two-in-one interactive segmentation tool was developed including a method based on the vanilla version of SAM2, and another based on the best classical method from the non-interactive evaluation in the previous section (see Section 3.1), to allow for a direct comparison in terms of tool efficiency. The SAM2 implementation was chosen due to its theoretical SOTA superiority, as the non-interactive evaluation described in the previous sectioncould not be used to conclude how they would perform when used interactively.

The following subsections go into details on how this tool was developed, the experimental design for the evaluation of interaction efficiency, and the definition of interaction metrics.

3.2.1 Tool architecture

The tool was developed as a web-application using Svelte as the frontend framework, and FastAPI as the backend framework. Due to it being a web application, it could be run remotely on any device with a web browser, without the need for installation, making it more easily accessible for online user testing with Ester Salgarella. And to make the development environment be consistent with production, the application was containerized using Docker, which also made it easier to deploy and run on any machine without worrying about dependencies and environment setup.

Svelte was chosen for the frontend mostly due to its reactive state management, . FASTAPI was chosen as the backend, as it is asynchronous, and a Python model-serving

[18] Rainio, Teuho, and Klén 2024.

3.2.1.1 Requirement analysis and first usability test

The 5th of May a requirement analysis and first usability test of a MVP was performed with Ester through a video call. That MVP was composed of two tools: a classical method (Gaussian), and an interactive SAM version (MobileSAM). After trying out the tools, she gave the impression that she could probably use the outputs as draft drawings to be completed with a pen, which might save her some time.

For the classical baseline, the grainy resolution (due to the input image quality) bothered her a bit. The output mask was also noisy at times. And not all lines were always picked up.

For the SAM version, she pointed out that it did not pick up the noisy tablet damage as much as the classical method did. But she was worried that the interaction with SAM would take longer than simply doing strokes manually. The box-prompt workflow was not completely successful during the test. Due to bugs, boxes mostly did not work. and when they did, the segmented forms were not useful.

3.2.1.2 Second (small) usability test

The 8th of May another usability test was performed on an updated MVP with Ester through a video call. Now, the box prompts worked, and the MobileSAM was replaced with SAM2.1. The segmentation quality per box-prompted sign was noticeably better than previously. Together with Lukas Esterle, the conclusions were to make a slider for the classical model to decrease and increase granularity. It was also suggested that box prompts for SAM could be computationally preselected based on the mask from the classical model, but this was discarded as it was not deemed needed.

3.2.1.3 Third usability and final acceptance test

The 12th of May, a third usability test was conducted in person at AIAS in Århus, including the first part of the usability experiment. As all 45 images could not be drawn all at once, the .ora exports from the tool were saved, and finished at a later time in batches by Ester on her own, for each of them recording the time it took to do so. Due to the huge

amount of time it took to process the images with the tool (especially with the SAM version), the experiment was stopped after just 3 images.

3.2.2 Experimental design for automation evaluation

3.2.3 Definition of interaction metrics

3.3 Workflow efficiency

To answer the question of RQ3: *How much reduction in annotation time can be achieved using the proposed tool compared to fully manual digitization, under controlled experimental conditions?*, a case study of 3 documents was conducted with Ester Salgarella, where the time taken to finish the segmentation outputs exported from the tool was measured, and compared to the time taken to do the same drawings fully manually.

The following subsections go into more details on the precautions taken in an attempt to fit the tool into the workflow, the experimental design for the evaluation of workflow efficiency, and the definition of workflow efficiency metrics.

3.3.1 Export to Krita

When the segmentation is complete, the tool allows for export to JPEG, PNG, or OpenRaster (.ora) format. The latter is an open file format that like Krita (.kra) files consists of a ZIP file with an XML document describing the layer structure, and a PNG or SVG image for each raster or vector layer, respectively. This is simpler in structure and easier to read and write for external programs than the .kra format, as it was specificially designed as an open exchange format between graphics editors. It is supported by a few image editing programs, including Krita, which makes it ideal for the exported file to be directly imported into Krita for further manual editing and annotation just by opening the file.

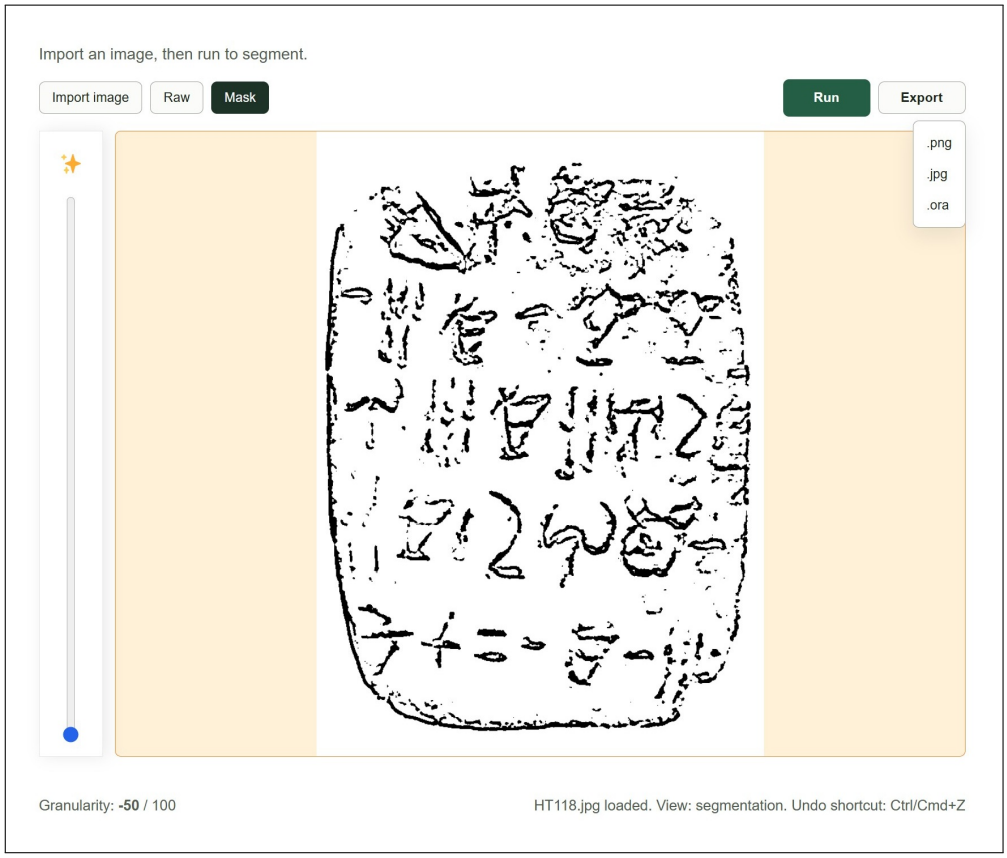

Figure 3.11: Yeah.

3.3.2 Definition of workflow efficiency metrics

Time.

3.3.3 Experimental design

The drawings were made in Krita using a displayless pressure-sensitive Wacom drawing pad with a pen input. The process leading to tool export, we recorded there in person, while Ester finished the polishing in Krita by her own the days afterward, recording the time taken and saving the final drawings

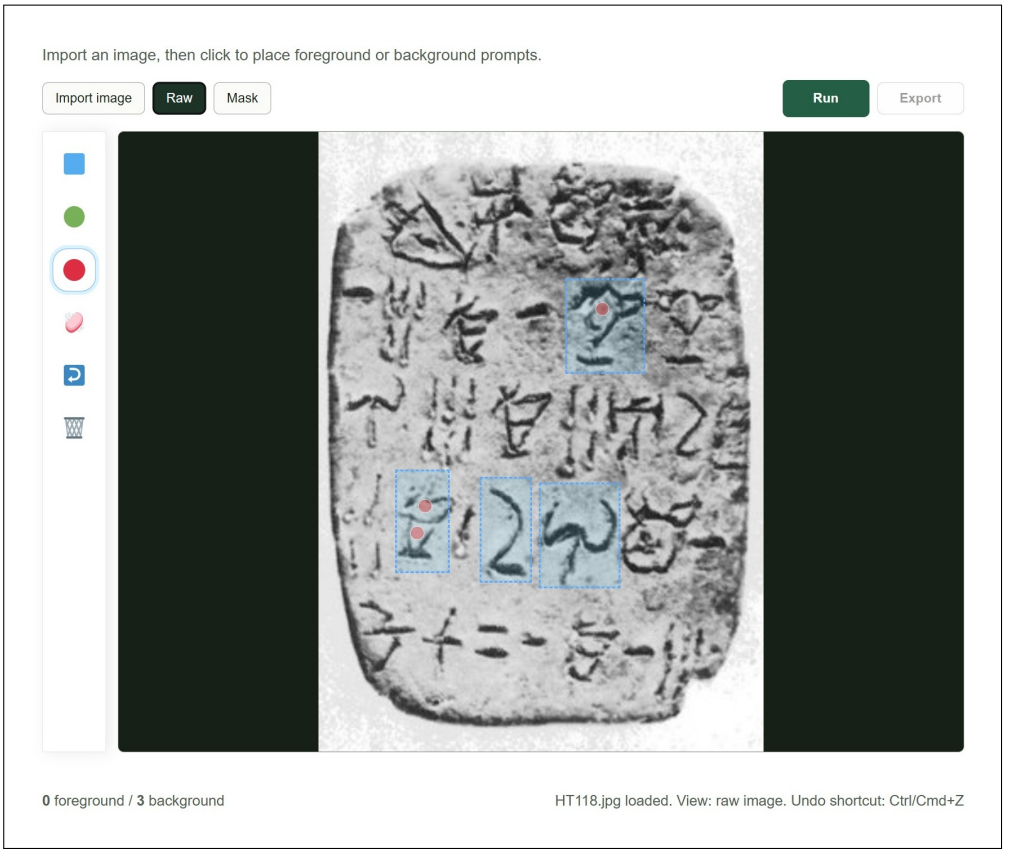

Figure 3.12: Yeah.

3.3.4 Analysis of results

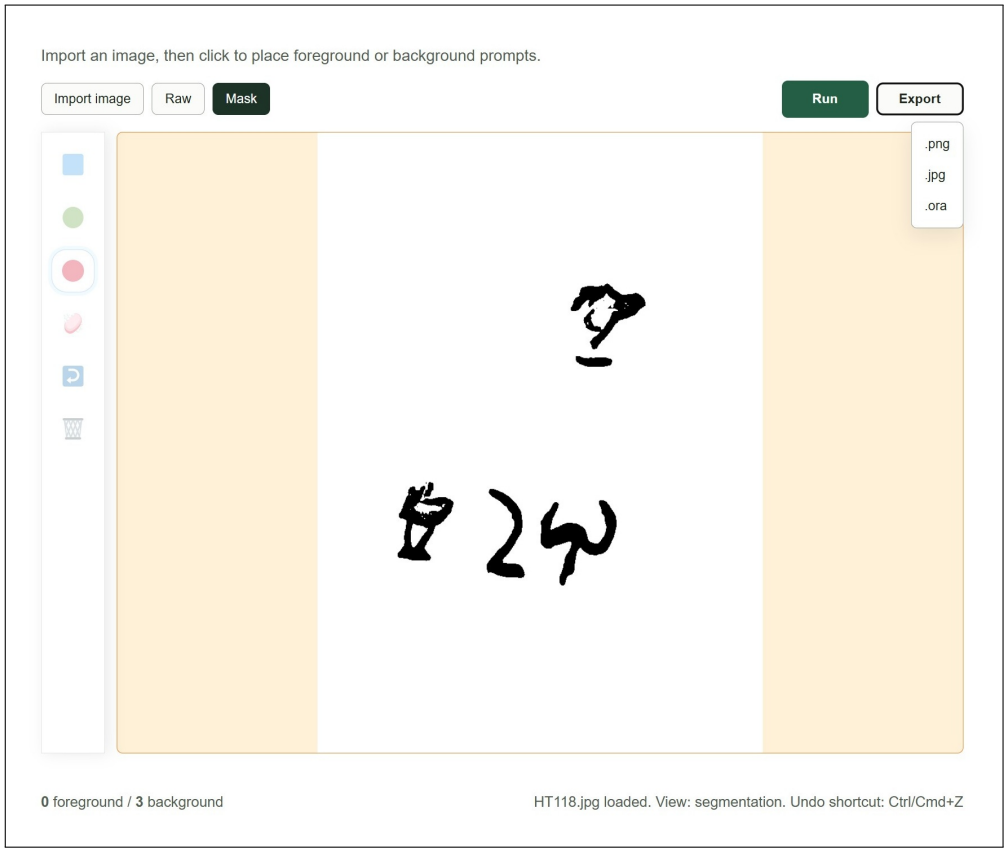

Figure 3.13: Yeah.

CHAPTER 4

Results and Discussion

4.1 Non-interactive segmentation performance

After running the non-interactive segmentation evaluation of the models on the dataset (hyperparameter optimization results can be found in Section A.1.1), it was clear that **mSAM!** did not perform better than the classical baselines (see Figure 4.3 for a plot of the results).

Meanwhile, SAM+pts with the engineered point prompt clouds performed considerably better. With a mean Dice score of around 0.217, it performed better than Otsu and CannyFill (with mean Dice scores of 0.224 and 0.183, respectively), and comparably to the best performing classical baseline, Gaussian that yielded a Dice score of 0.327.

The following subsections go into a bit more detail on the explanation of these results, the statistical significance of the differences in performance between SAM and the best classical baseline (see Section A.1.2 in the appendix for the full results table), and a discussion of the fairness of this experiment.

4.1.1 SAM ablation

The results of the ablation study of different SAM variants (see Figure 4.1) showed that the best performing variant was the one with automatic point prompts and bilateral filter, making use of the single output token per image instead of best of three. This was also the one used for the main evaluation.

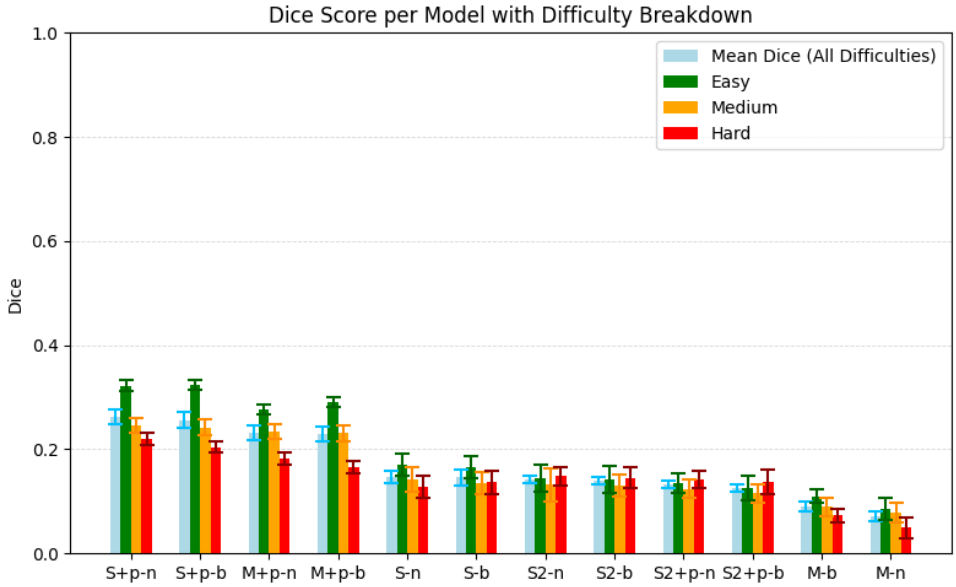

Figure 4.1: Dice Score performance per SAM variant. Split by difficulty. With error bars representing the standard error of the mean. Ordered descendingly from the left by mean performance.

4.1.2 Few-shot SAM

The few-shot adaptation of training free SAM (FATESAM) did not perform better than the default SAM implementation, which is not surprising given the small size and poor quality of the dataset, and poor alignment of support images with ground truth.

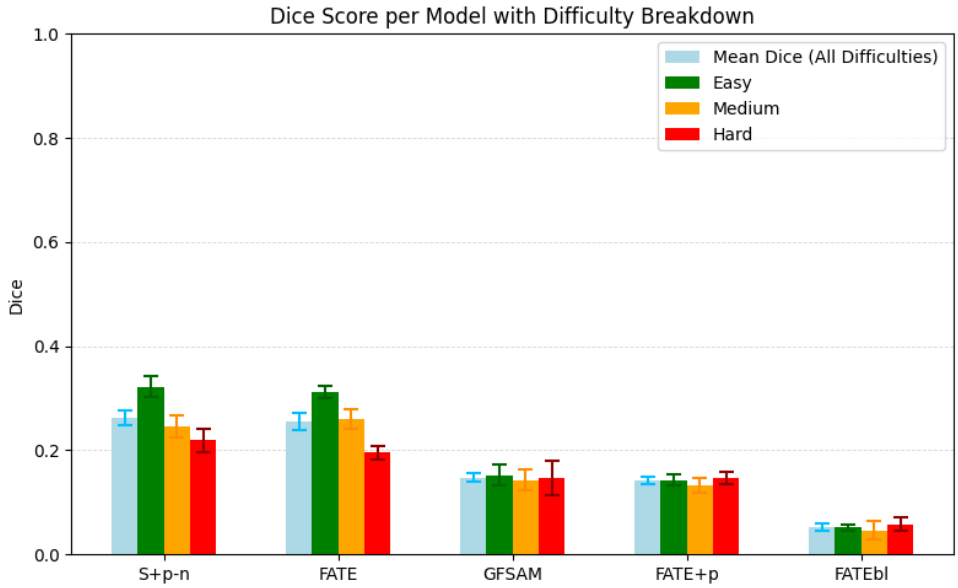

Figure 4.2: Dice score comparison of variants of FATESAM with the best performing SAM implementation.

4.1.3 Model comparison

Knowing the overall quantitative results (see Figure 4.3), it pays to look at a visual example to see what is going on behind the scenes.

CannyFill also oversegments, due to its large kernel window size for morphological closing, which was necessary to fill in the gaps between edges, but gives a blocky look, and merges all nearby edges together.

The unprompted SAM segments only background noise here, producing nothing of value. But the prompted SAM+pts does produce a somewhat coherent mask, capturing almost every sign with less oversegmented areas than both Otsu and CannyFill. Still, it oversegments many signs (especially those with background-space supposed to be inside of them), fades out signs where they are supposed to continue (due to loss of attention, as segmenting many separated signs at once is not what its attention weights have been trained for), and is not able to capture the outline of the tablet. But the signs that it does capture well, are captured with a good level of spatial coherence. And there is way less

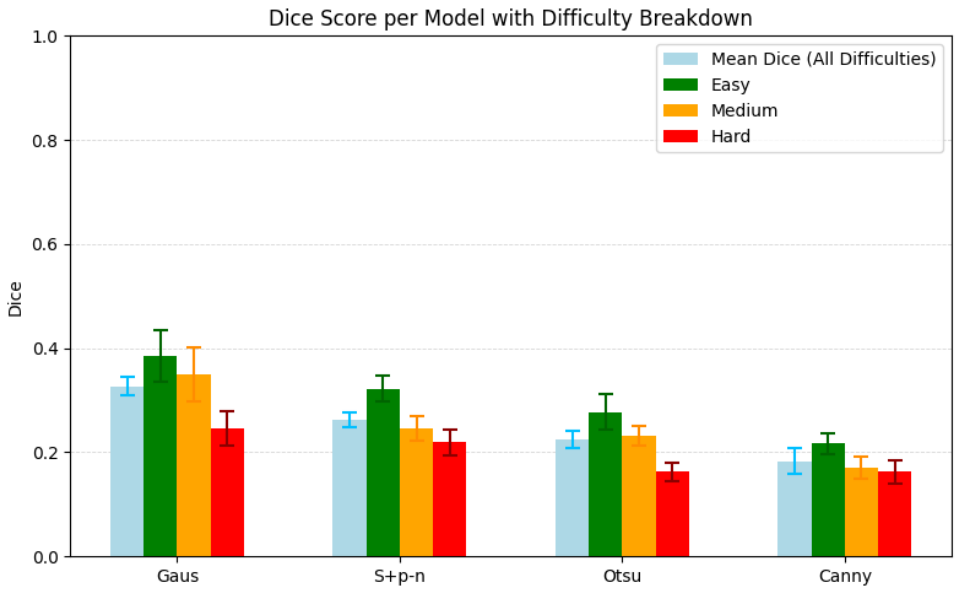

Figure 4.3: Dice score comparison of all models on the evaluation dataset.

noise between them compared to the masks of the baselines, which is arguably a positive sign for the potential of SAM to perform well on this task once properly fine-tuned prompts, or points are placed through human supervision.

((a)) Raw Image ((b)) Ground Truth ((c)) Otsu 0.292

((d)) Gaussian 0.445 ((e)) CannyFill 0.39 ((f)) SAMpts 0.38

Figure 4.4: Visual comparison of output masks per model from segmentation on HT7a (easy), shown with Dice scores next to the raw image and ground truth.

4.1.4 Statistical significance

As it was clear from the results, SAM+pts did not perform better than the best performing classical baseline. To determine whether it performed significantly worse, a paired t-test was conducted, as residuals of the differences in Dice scores between Gaussian and SAM+pts were approximately normally distributed (see next Section 4.1.4.1). What follows is the documentation of that normality test, the paired t-test, and the results of both.

4.1.4.1 Normality of residuals

As the QQplot in Figure 4.5 shows, the residuals of the paired differences in Dice scores between Gaussian and SAM+pts appeared to be approximately normally distributed. Furthermore, the Shapiro-Wilk test for normality of these residuals yielded a p-value of $0.056 > \alpha = 0.05$, indicating that the null hypothesis of normality could not be rejected. Thus, the normality assumption for the paired t-test was satisfied, and its results could be considered valid.

4.1.4.2 Paired t-test

The paired t-test was conducted on the Dice scores obtained for each of the 27 images across Gaussian and SAM+pts. The null hypothesis (H_0) stated that there is no difference in mean Dice scores between the two models, while the alternative hypothesis (H_1) stated that Gaussian has a higher mean Dice score than SAM+pts.

Performing the paired t-test yielded a test statistic of $t \approx 3.737$ and a p-value of $P(T > t) \approx 0.001$. Thus, as $0.001 < \alpha = 0.05$, the null hypothesis could be rejected, and it could be concluded that Gaussian significantly outperforms SAM+pts in terms of Dice score. Hence, it can be said that SAM performed significantly worse in comparison to the best classical segmentation methods for non-interactive segmentation of Linear A tablets, even when automatically prompted.

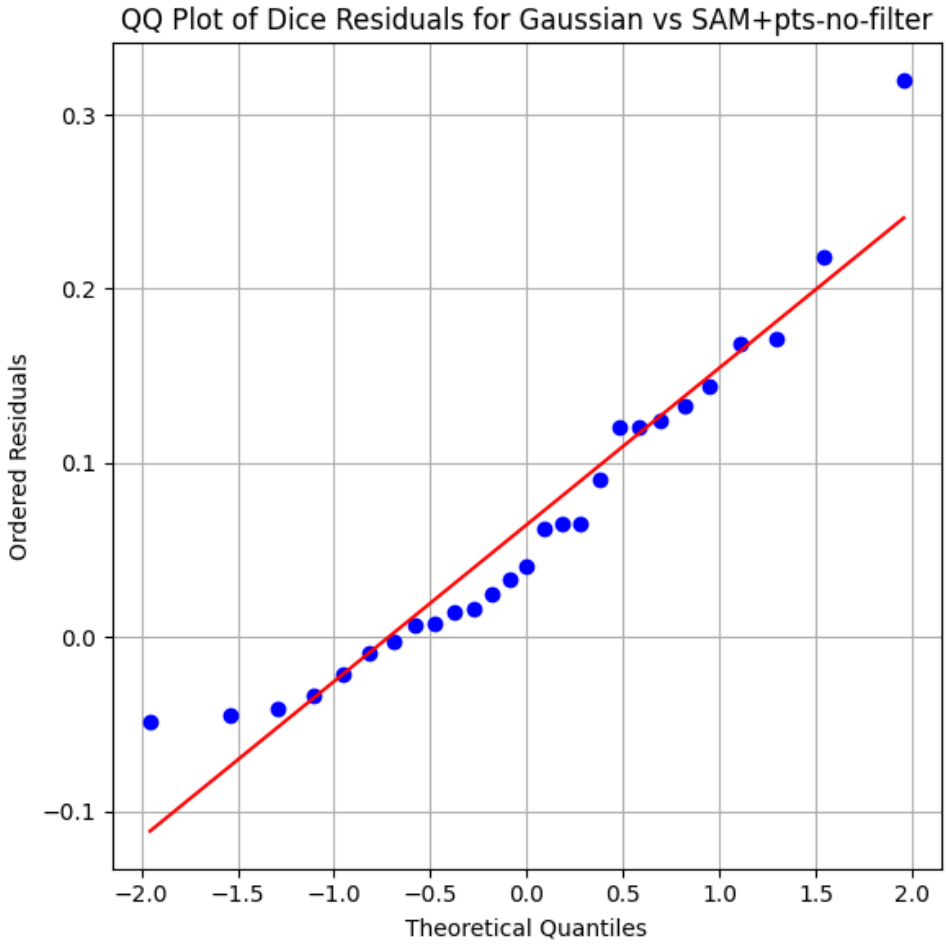

Figure 4.5: QQplot of the residuals of the paired t-test comparing Gaussian and SAM+pts. The residuals appear to be approximately normally distributed, which supports the validity of the paired t-test results. Furthermore, the Shapiro-Wilk test for normality of these residuals yielded a p-value of $0.056 > \alpha = 0.05$, indicating that the null hypothesis of normality could not be rejected.

4.1.5 Fairness of the experiment

4.1.5.1 Data

The dataset used for evaluation was the same as the one used for hyperparameter optimization of models, which could be seen as a potential source of bias. Additionally, the dataset of 60 images is quite small, which makes it more susceptible to noise and outliers, although sufficiently sized to support the statistical analysis. Moreover, the ground truth annotations of the dataset were heavily biased, even as they were manually created and registered to the raw images. Therefore, Dice scores obtained from the evaluation were unreasonably low, and the performance of the models is not properly reflected by these scores. However, for comparison between models, the relative differences in Dice scores were still meaningful, as all models were evaluated on the same imperfectly aligned dataset.

4.1.5.2 Model implementation

In the case of the fairness of the SAM implementation chosen, the experiment obviously only evaluates the performance of an untuned version. This implementation had most certainly never seen Linear A tablet images during training, as it was pretrained by default on the large and diverse dataset of natural images SA-1B. Furthermore, the model implementation used makes use of MobileSAM with its small ViT image encoder TinyVit (Wu et al. 2022), which can compromise on segmentation quality. Lastly, SAM was restricted by the automatic nature of this experiment, as its intended use is interactive segmentation. Thus, it does not come as a surprise that **mSAM!** underperforms in comparison to the classical baselines. The results therefore suggest a need for human intervention for the optimal placement of prompts to provide the best conditions for SAM, as it otherwise underperforms.

4.2 Tool-side segmentation performance

The case study of 3 documents showed that the interactive version of SAM2 with box prompts and point prompts not only took more time to use than the classical baselinedue to its the burden of interaction, but also did not perform better in terms of segmentation quality, even with the help of human input (see Figure 4.6).

4.3 Workflow impact

The case study of 3 documents

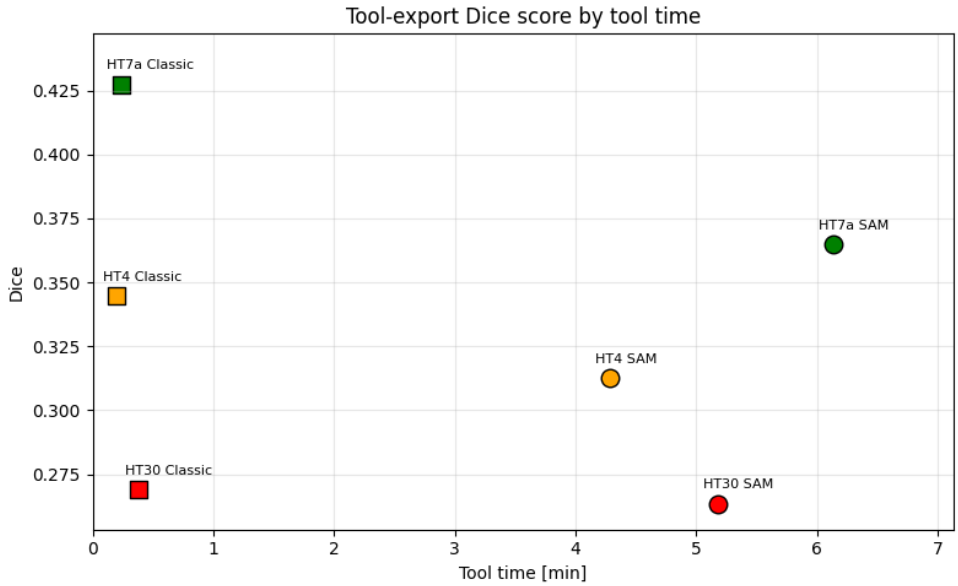

Figure 4.6: Dice score performance of the interactive version of SAM2, compared to the classical baselines, plotted against the time taken to use each tool.

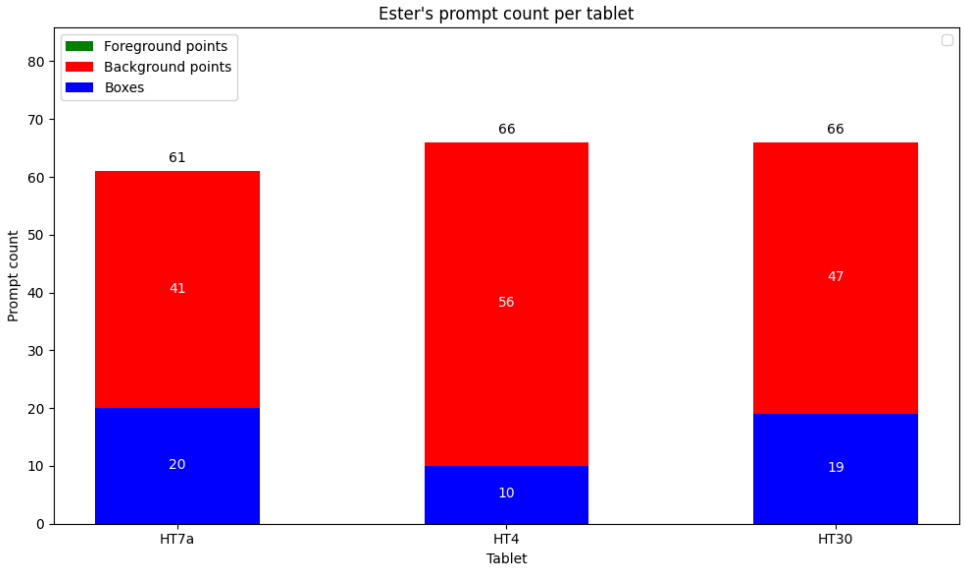

Figure 4.7: Amount of prompts placed per image for the interactive version of SAM2.

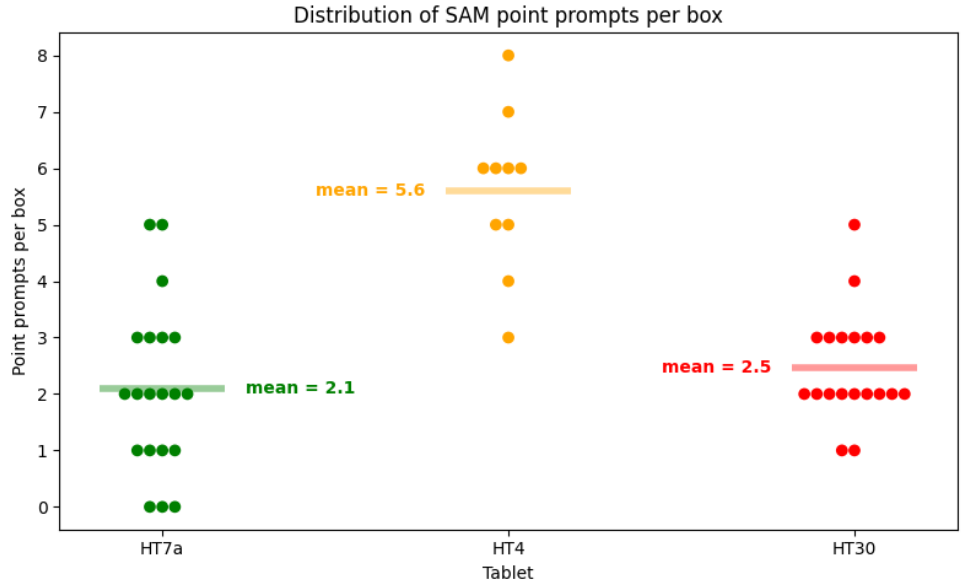

Figure 4.8: Average number of point prompts placed per box prompt per image for the interactive version of SAM2.

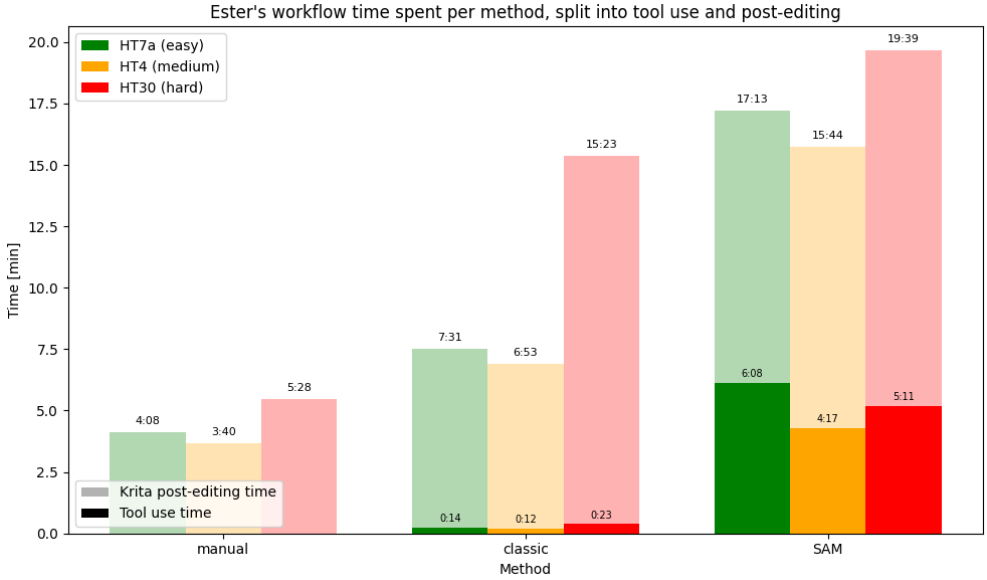

Figure 4.9: Workflow time spent per method, split into tool use

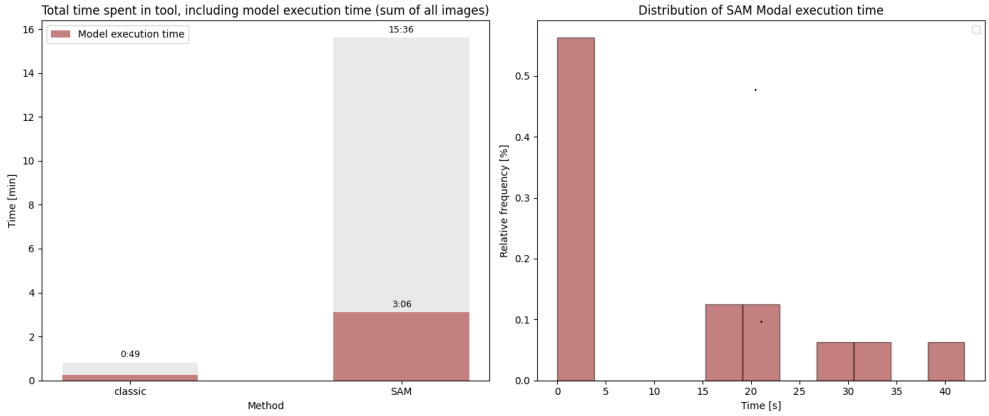

Figure 4.10: This is how to do it

CHAPTER 5

Conclusion

5.1 Summary of findings per research question

5.1.1 RQ1: Segmentation performance

5.1.2 RQ2: Degree of automation

5.1.3 RQ3: Workflow efficiency

5.2 Overall contribution

5.3 Future Work

5.3.1 Methodological extensions

5.3.2 Dataset expansion and diversification

5.3.3 Further automation opportunities

...The results therefore suggest a need for a generalized and computationally efficient SAM implementation to be tuned to the greater domain of inscriptions, as the default model is not sufficient enough to prove its potential capabilities there. For this potential tuning to be possible,

there is a need of a sufficiently sized one-class-labeled dataset of general inscription segmentations.

Bibliography

Bergstra, James, Remi Bardenet, and Balázs Kégl (December 2011). "Algorithms for Hyper-Parameter Optimization." In.

Bergstra, James, Daniel Yamins, and David D. Cox (2013). "Making a Science of Model Search: Hyperparameter Optimization in Hundreds of Dimensions for Vision Architectures." In: *Proceedings of the 30th International Conference on Machine Learning (ICML)*. Volume 28. Proceedings of Machine Learning Research 1. PMLR, pages 115–123. URL: https://proceedings.mlr.press/v28/bergstra13.pdf.

Carion, Nicolas et al. (2025). *SAM 3: Segment Anything with Concepts*. arXiv: 2511.16719 [cs.CV]. URL: https://arxiv.org/abs/2511.16719.

Christoffersen, Frederik W. L. (2026). *Source code for this Thesis*. Source code repository. URL: https://github.com/FrederikWLC/linA-scribe.

Consiglio Nazionale delle Ricerche (2025). *A Major Addition to the Linear A Corpus*. URL: https://www.cnr.it/en/news/13529/a-major-addition-to-the-linear-a-corpus (visited on February 4, 2026).

Del Freo, Maurizio and Julien Zurbach (2025). *Recueil des inscriptions en linéaire A. Supplément 1 (RILA-S1)*. Athens: École française d'Athènes.

Dosovitskiy, Alexey et al. (2021). *An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale*. arXiv: 2010.11929 [cs.CV]. URL: https://arxiv.org/abs/2010.11929.

Godart, Louis and Jean-Pierre Olivier, editors (1976-1985). *Études crétoises. Recueil des inscriptions en linéaire A (GORILA)*. Five-volume corpus of Linear A inscriptions (GORILA I–V). Athènes. URL: https://cefael.efa.gr/result.php?site_id=1&serie_id=EtCret.

Gonzalez, Rafael C. and Richard E. Woods (2018). *Digital Image Processing*. 4th edition. New York: Pearson. ISBN: 978-0133356724. URL:

https://www.cl72.org/090imagePLib/books/Gonzales,Woods-Digital.Image.Processing.4th.Edition.pdf.

Hamplová, Adéla et al. (2024). "Instance Segmentation of Characters Recognized in Palmyrene Aramaic Inscriptions." In: *Computer Modeling in Engineering & Sciences* 140.3, pages 2869–2889. ISSN: 1526-1506. DOI: 10.32604/cmes.2024.050791. URL: http://www.techscience.com/CMES/v140n3/57249.

Han, Dongsheng et al. (2023). *Segment Anything Model (SAM) Meets Glass: Mirror and Transparent Objects Cannot Be Easily Detected*. arXiv: 2305.00278 [cs.CV]. URL: https://arxiv.org/abs/2305.00278.

He, Xingxin et al. (2025). *Few-Shot Adaptation of Training-Free Foundation Model for 3D Medical Image Segmentation*. arXiv: 2501.09138 [cs.CV]. URL: https://arxiv.org/abs/2501.09138.

Kirillov, Alexander et al. (2023). "Segment Anything." In: *arXiv preprint arXiv:2304.02643*. URL: https://arxiv.org/pdf/2304.02643.

Ma, Jun et al. (January 2024). "Segment anything in medical images." In: *Nature Communications* 15.1. ISSN: 2041-1723. DOI: 10.1038/s41467-024-44824-z. URL: http://dx.doi.org/10.1038/s41467-024-44824-z.

Modal Labs (2026). *Modal Documentation*. Documentation for the Modal cloud execution platform. URL: https://modal.com/docs/guide (visited on May 23, 2026).

Mohammadi, Mobin, Kaveh Mollazade, and Nasser Behroozi-Khazaei (2025). "Under- and over-segmentation: New metrics for image segmentation accuracy measurement." In: *Array* 28, page 100624. ISSN: 2590-0056. DOI: https://doi.org/10.1016/j.array.2025.100624. URL: https://www.sciencedirect.com/science/article/pii/S2590005625002516.

Oliveira, Saïd, Benoît Seguin, and Frédéric Kaplan (2018). "dhSegment: A generic deep-learning approach for document segmentation." In: *arXiv preprint arXiv:1804.10371*. URL: https://arxiv.org/abs/1804.10371.

OpenCV (2024). *cv::AdaptiveThresholdTypes — OpenCV Documentation*. OpenCV. URL: https://docs.opencv.org/4.x/d7/d1b/group__imgproc__misc.html#gaa42a3e6ef26247da787bf34030ed772c (visited on March 7, 2026).

Otsu, Nobuyuki (1979). "A Threshold Selection Method from Gray-Level Histograms." In: *IEEE Transactions on Systems, Man, and Cybernetics* 9.1, pages 62–66. DOI: 10.1109/TSMC.1979.4310076. URL: https://engineering.purdue.edu/kak/computervision/ECE661.08/OTSU_paper.pdf.

Pan, Cong (2025). "Real-Time Hardware Architecture Design for An Innovative Bilateral Filtering Algorithm." In: *Proceedings of the 2025 5th International Conference on Automation Control, Algorithm and Intelligent Bionics (ACAIB '25)*. New York, NY, USA: ACM, pages 102–107. ISBN: 9798400714313. DOI: 10.1145/3760269.3760285. URL: https://dl.acm.org/doi/10.1145/3760269.3760285 (visited on March 23, 2026).

Price, Stephen, Elke Rundensteiner, and Danielle L. Cote (2025). "Mastering SAM Prompts: A Large-Scale Empirical Study in Segmentation Refinement for Scientific Imaging." In: *Transactions on Machine Learning Research*. ISSN: 2835-8856. URL: https://openreview.net/forum?id=cWcTQMpqv6.

Rainio, Oona, Jarmo Teuho, and Riku Klén (2024). "Evaluation metrics and statistical tests for machine learning." In: *Scientific Reports* 14, page 6086. DOI: 10.1038/s41598-024-56706-x. URL: https://doi.org/10.1038/s41598-024-56706-x.

Ravi, Nikhila et al. (2024). *SAM 2: Segment Anything in Images and Videos*. arXiv: 2408.00714 [cs.CV]. URL: https://arxiv.org/abs/2408.00714.

Rosebrock, Adrian (2015). *Zero-parameter, automatic Canny edge detection with Python and OpenCV*. Accessed: 2026-04-06. URL: https://pyimagesearch.com/2015/04/06/zero-parameter-automatic-canny-edge-detection-with-python-and-opencv/.

Ryali, Chaitanya et al. (2023). *Hiera: A Hierarchical Vision Transformer without the Bells-and-Whistles*. arXiv: 2306.00989 [cs.CV]. URL: https://arxiv.org/abs/2306.00989.

Salgarella, Ester (2020). *Aegean Linear Script(s): Rethinking the Relationship Between Linear A and Linear B*. Cambridge: Cambridge University Press.

— (2025). *Writing in Bronze Age Crete: 'Minoan' Linear A*. Cambridge: Cambridge University Press.

Salgarella, Ester and Simon Castellan (2020). *SigLA: The Signs of Linear A: a Palaeographical Database*. Research Paper / Technical Report. Accessed: 2026-02-04. SigLA. URL: https://sigla.phis.me/paper.pdf.

SigLA (2021a). *About SigLA*. URL: https://sigla.phis.me/about.html (visited on February 4, 2026).

— (2021b). *SigLA Document Search Interface*. URL: https://sigla.phis.me/search-document.html (visited on February 4, 2026).

Tang, Lv, Haoke Xiao, and Bo Li (2023). *Can SAM Segment Anything? When SAM Meets Camouflaged Object Detection*. arXiv: 2304.04709 [cs.CV]. URL: https://arxiv.org/abs/2304.04709.

Tomasi, Carlo and Roberto Manduchi (1998). "Bilateral Filtering for Gray and Color Images." In: *Proceedings of the Sixth International Conference on Computer Vision*. Bombay, India: IEEE, pages 839–846. DOI: 10.1109/ICCV.1998.710815. URL: https://users.soe.ucsc.edu/~manduchi/Papers/ICCV98.pdf (visited on March 23, 2026).

Vaswani, Ashish et al. (2023). *Attention Is All You Need*. arXiv: 1706.03762 [cs.CL]. URL: https://arxiv.org/abs/1706.03762.

Wang, Shibin et al. (2025). "OBIFlo-SAM: multi-task semantic recognition and segmentation of oracle bone inscription." In: *npj Heritage Science* 13, page 588. DOI: 10.1038/s40494-025-02157-0. URL: https://www.nature.com/articles/s40494-025-02157-0.

Wu, Kan et al. (2022). "TinyViT: Fast Pretraining Distillation for Small Vision Transformers." In: *arXiv preprint arXiv:2207.10666*. arXiv: 2207.10666 [cs.CV]. URL: https://arxiv.org/abs/2207.10666.

Ye, Maoyuan et al. (2024). *Hi-SAM: Marrying Segment Anything Model for Hierarchical Text Segmentation*. arXiv: 2401.17904 [cs.CV]. URL: https://arxiv.org/abs/2401.17904.

Zhang, Chaoning et al. (2023). "Faster Segment Anything: Towards Lightweight SAM for Mobile Applications." In: *arXiv preprint arXiv:2306.14289*. URL: https://arxiv.org/pdf/2306.14289.

Zhang, Pan, Chao Li, and Yuanhua Sun (2024). "Stone Inscription Image Segmentation Based on Stacked-UNets and GANs." In: *Discover Applied Sciences* 6.10, page 550. DOI: 10.1007/s42452-024-06264-8. URL: https://doi.org/10.1007/s42452-024-06264-8.

Zhao, Xu et al. (2023). "MobileSAMv2: Faster Segment Anything to Everything." In: *arXiv preprint arXiv:2312.09579*. URL: https://arxiv.org/pdf/2312.09579.

APPENDIX A

An Appendix

A.1 Non-interactive segmentation results

A.1.1 Hyperparameter optimization

Table A.1: Best TPE hyperparameter optimization results for CannyFill after 100 trials.

Hyperparameter	Obtained value
$d_{bilateral}$	23
$\sigma_{bilateral}$	21
$d_{closing}$	19
σ_{canny}	$8.7 \cdot 10^{-1}$

Table A.2: Best TPE hyperparameter optimization results for Gaussian after 100 trials.

Hyperparameter	Obtained value
$d_{bilateral}$	25
$\sigma_{bilateral}$	43
C	9
$d_{gaussian}$	39

Table A.3: Best TPE hyperparameter optimization results for SAM+pts after 100 trials.

Hyperparameter	Obtained value
$d_{bilateral}$	18
σ	141
C	5
$d_{gaussian}$	23
n_{fgd_pts}	250
n_{bgd_pts}	1352
$d_{gap\_erosion}$	5

A.1.2 Raw evaluation data

Table A.4: Raw Dice score data on easy images. Label column represents the name of the document segmented, while all other columns represent Dice scores obtained for those documents for a given model.

Label	CannyFill	Gaussian	Otsu	SAM+pts-no-filter
HT10a	0.39	0.445	0.292	0.38
HT131a	0.152	0.464	0.401	0.343
HT17	0.253	0.367	0.157	0.389
HT21	0.191	0.402	0.22	0.278
HT22	0.226	0.382	0.319	0.32
HT7a	0.403	0.427	0.415	0.259
MA9	0.001	0.204	0.107	0.213
PH2	0.324	0.423	0.245	0.383
PH31a	0.009	0.354	0.343	0.338

Table A.5: Raw Dice score data on medium images. Label column represents the name of the document segmented, while all other columns represent Dice scores obtained for those documents for a given model.

Label	CannyFill	Gaussian	Otsu	SAM+pts-no-filter
HT13	0.422	0.451	0.301	0.233
HT132	0.001	0.472	0.208	0.153
HT135a	0.107	0.343	0.249	0.223
HT154B	0	0.303	0.186	0.27
HT20	0.181	0.353	0.177	0.209
HT4	0.305	0.356	0.321	0.348
HT85b	0.317	0.293	0.231	0.287
HT90	0.188	0.266	0.251	0.311
PH31b	0.014	0.313	0.164	0.18

Table A.6: Raw Dice score data on hard images. Label column represents the name of the document segmented, while all other columns represent Dice scores obtained for those documents for a given model.

Label	CannyFill	Gaussian	Otsu	SAM+pts-no-filter
HT1	0.131	0.181	0.096	0.229
HT115b	0.059	0.19	0.09	0.124
HT122b	0.212	0.268	0.138	0.177
HT140	0.225	0.237	0.158	0.213
HT3	0.141	0.155	0.133	0.158
HT30	0.282	0.271	0.232	0.312
HT8a	0.295	0.402	0.182	0.231
HT9b	0.1	0.212	0.205	0.197
KN1a	0.02	0.3	0.228	0.334

A.1.3 Summary statistics

Table A.7: Summary Statistics per Model of Dice Score on all images.

Model	Mean Dice	Standard deviation	Standard error	n
CannyFill	0.183	0.132	0.025	27
Gaussian	0.327	0.093	0.018	27
Otsu	0.224	0.087	0.017	27
SAM+pts-no-filter	0.263	0.076	0.015	27

Table A.8: Summary Statistics per Model of Dice Score on easy images.

Model	Mean Dice	Standard deviation	Standard error	n
CannyFill	0.217	0.147	0.049	9
Gaussian	0.385	0.077	0.026	9
Otsu	0.278	0.105	0.035	9
SAM+pts-no-filter	0.322	0.061	0.02	9

Table A.9: Summary Statistics per Model of Dice Score on medium images.

Model	Mean Dice	Standard deviation	Standard error	n
CannyFill	0.171	0.154	0.051	9
Gaussian	0.35	0.07	0.023	9
Otsu	0.232	0.054	0.018	9
SAM+pts-no-filter	0.246	0.063	0.021	9

Table A.10: Summary Statistics per Model of Dice Score on hard images.

Model	Mean Dice	Standard deviation	Standard error	n
CannyFill	0.163	0.097	0.032	9
Gaussian	0.246	0.075	0.025	9
Otsu	0.162	0.053	0.018	9
SAM+pts-no-filter	0.22	0.068	0.023	9

$\sqrt{17}$
$\Omega \int \delta e^{i\pi} = -1$
{2.7182818284}ϑφερτυθιοπσδφγηξκλ
χ^2
$\sum$
λ

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Technical
University of
Denmark

Richard Petersens Plads, Building 324
2800 Kgs. Lyngby
Tlf. 4525 1700

www.compute.dtu.dk